



Hunting Performance in Python Code

twitter.com/sumercip

github.com/sumerc

sumercip.com

Outline

- Why?
- How?
- What?
- A slight peek over the ecosystem
- Some common rules for better performance
- Q/A



Performance = Measurement



1,343 repository results

Sort: Best match ▾

 [rkern/line_profiler](#)

Line-by-line **profiling** for **Python**

★ 3.3k ● Python Updated on 11 Dec 2019

 [pythonprofilers/memory_profiler](#)

Monitor Memory usage of **Python** code

★ 2.1k ● Python Updated on 13 Jan

 [nvdv/vprof](#)

Visual **profiler** for **Python**

[cpu-flame-graph](#) [stats](#) [python](#) [visualization](#) [javascript](#) [d3](#) [profiler](#) [developer-tools](#)

★ 3.6k ● Python BSD-2-Clause license Updated on 29 Dec 2019

 [benfred/py-spy](#)

Sampling **profiler** for **Python** programs

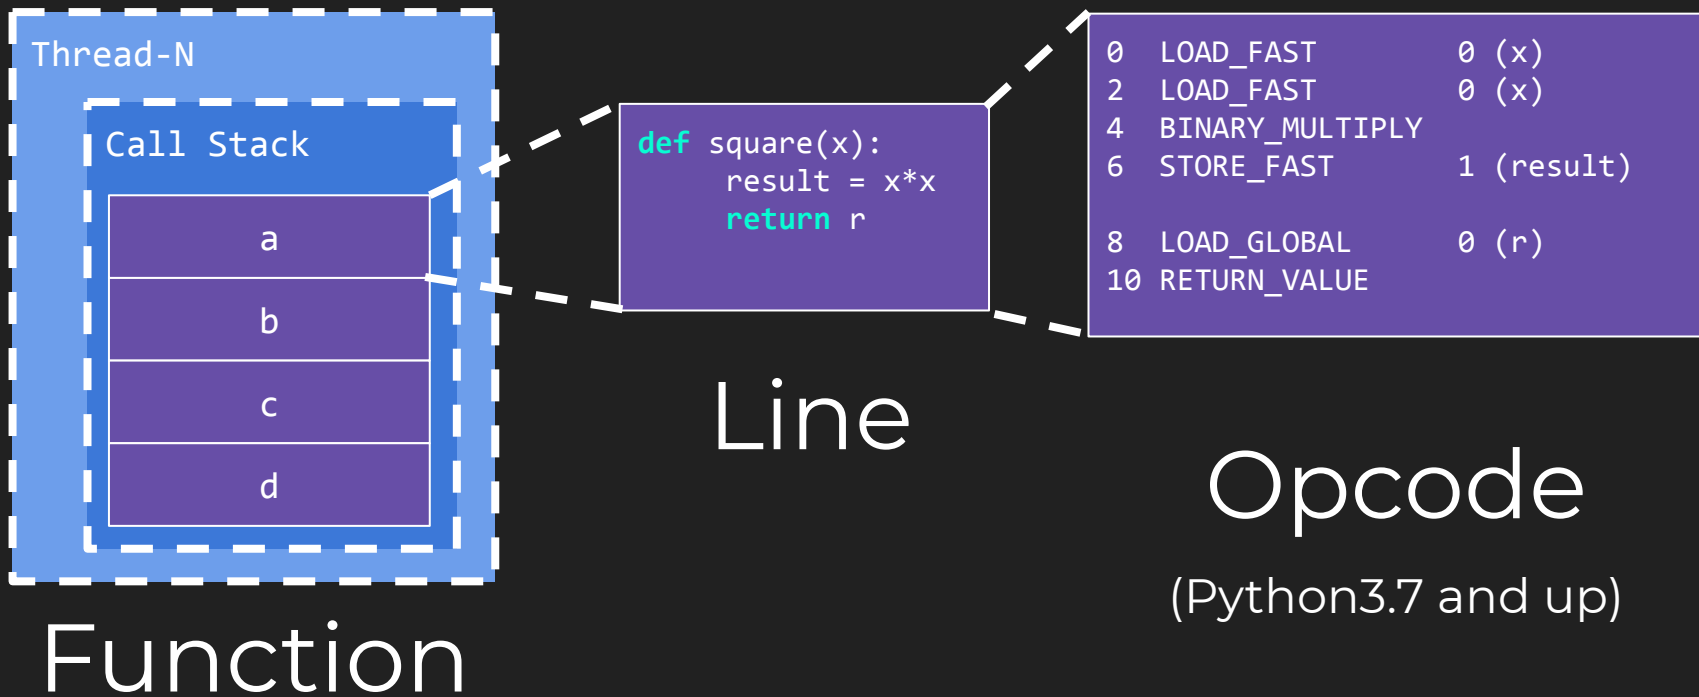
[python](#) [profiler](#) [performance-analysis](#) [profiling](#)

★ 5.1k ● Rust MIT license Updated 4 days ago

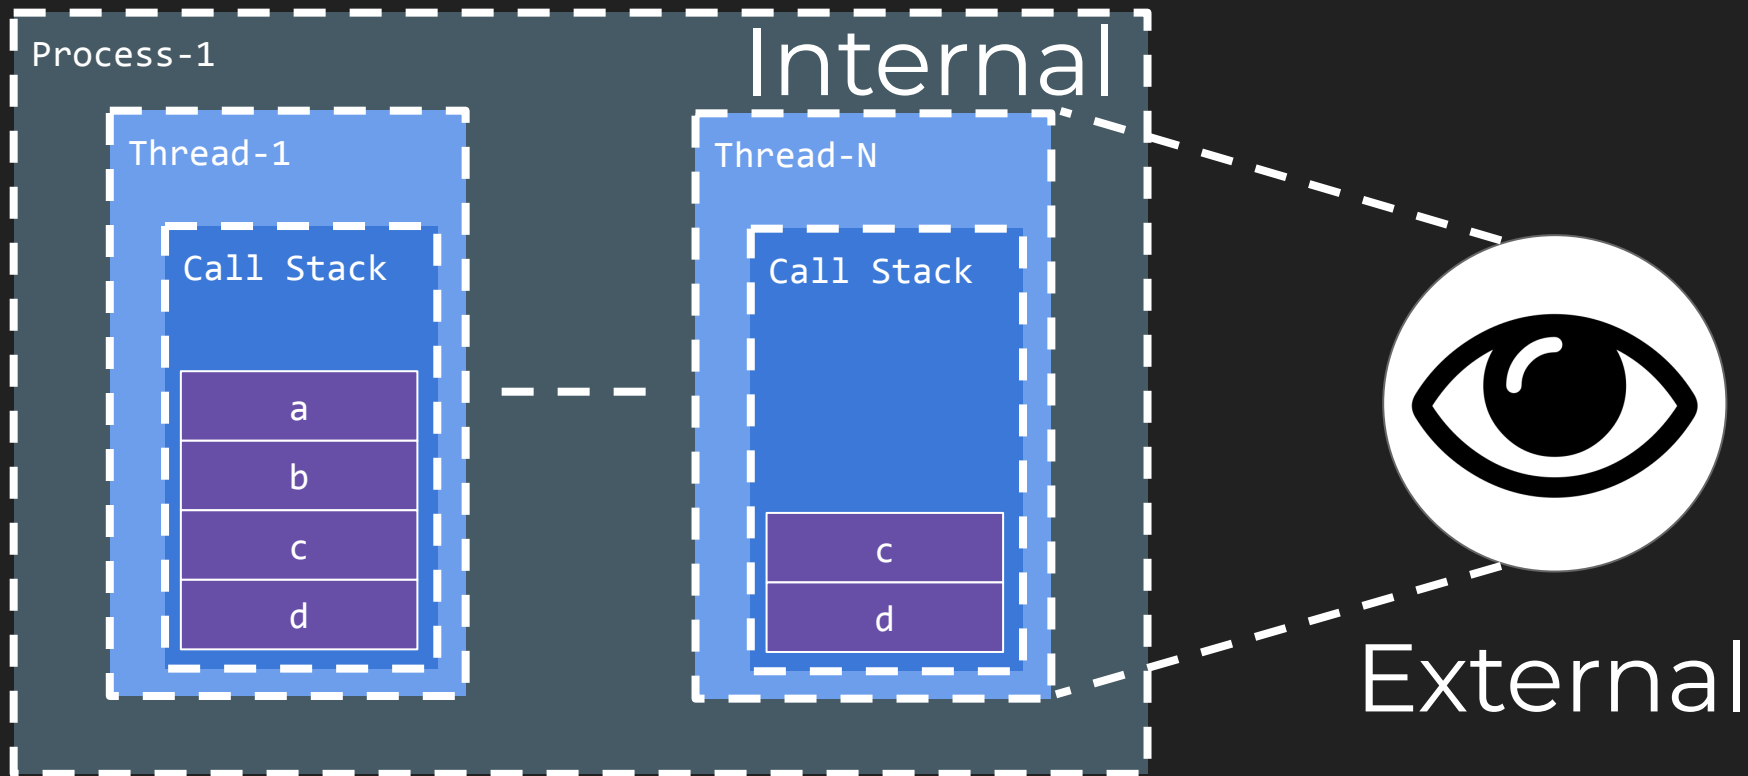


Tracing vs Sampling

TRACING



SAMPLING



What does profilers measure?



Time?
Yes. But which time?

Wall

CPU

I/O (Wall - CPU)

More metrics?

Memory

tracemalloc, objgraph

App. specific metrics

django-debug-toolbar

What does profilers output?

Outputs

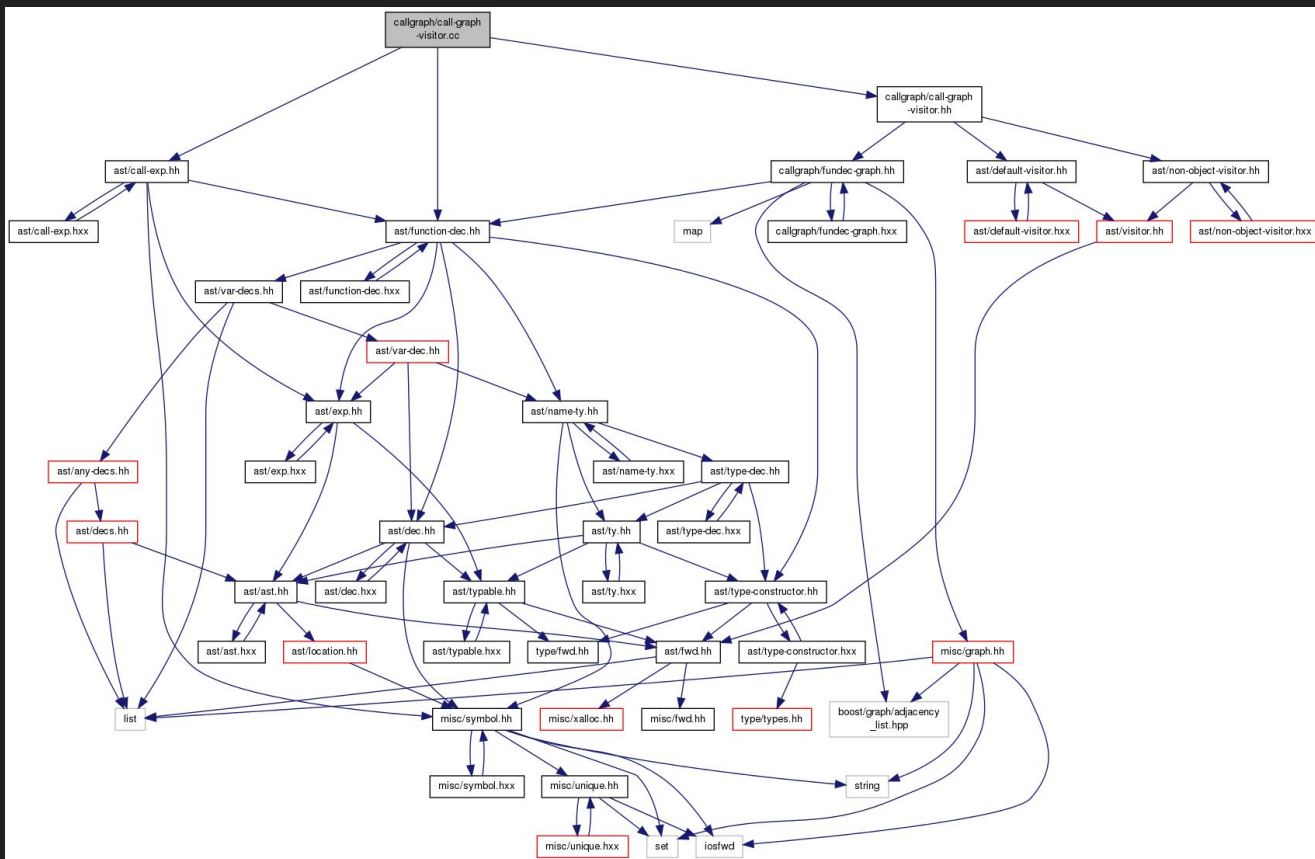
- Table
- Callgraph
- Flamegraph
- Custom formats

Table

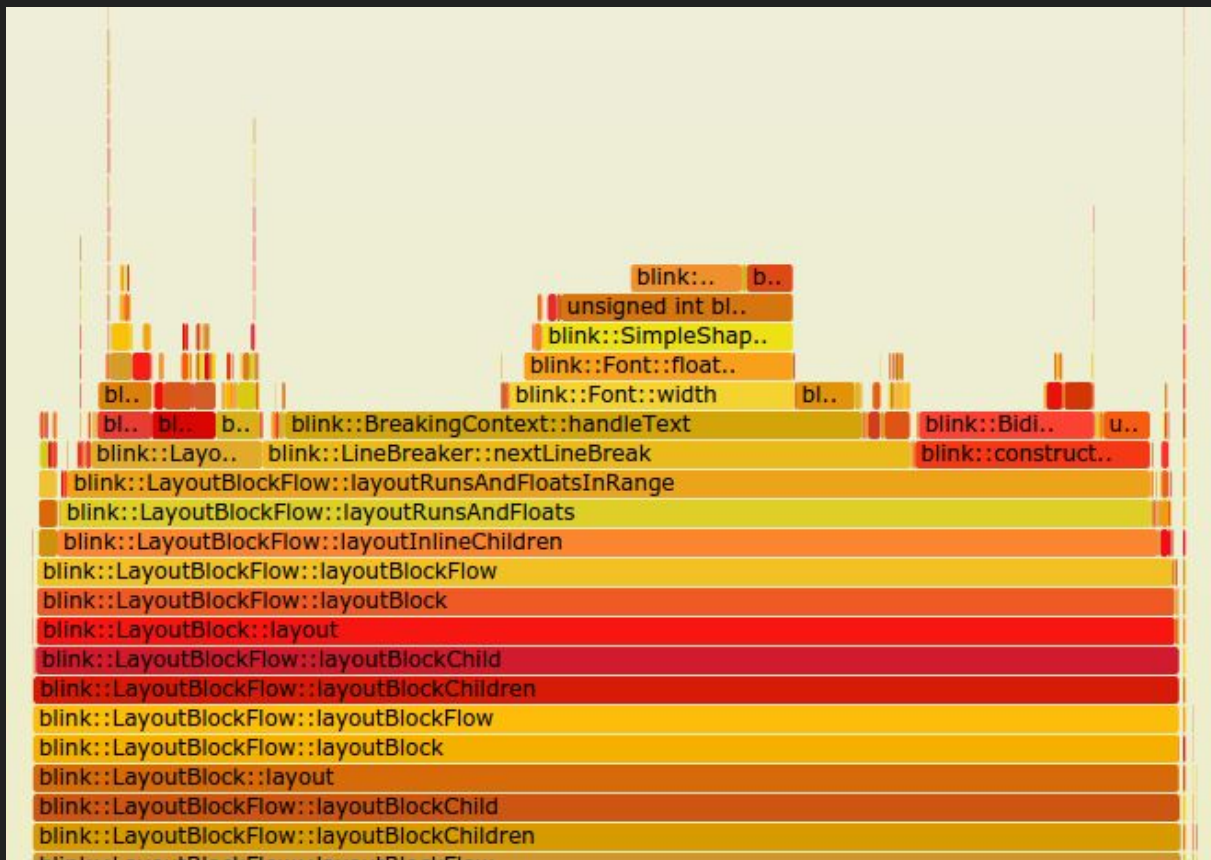
Ordered by: cumulative time

ncalls	tottime	percall	cumtime	percall	filename:lineno(function)
2/1	0.000	0.000	0.002	0.002	{built-in method builtins.exec}
1	0.000	0.000	0.002	0.002	cprofile_demo.py:1(<module>)
1	0.001	0.001	0.001	0.001	cprofile_demo.py:3(test_multiply)
1	0.000	0.000	0.001	0.001	<frozen importlib._bootstrap>:966(_find_and_load)
1	0.000	0.000	0.001	0.001	<frozen importlib._bootstrap>:936(_find_and_load_unlocked)
1	0.000	0.000	0.001	0.001	<frozen importlib._bootstrap>:651(_load_unlocked)
1	0.000	0.000	0.000	0.000	<frozen importlib._bootstrap_external>:672(exec_module)
1	0.000	0.000	0.000	0.000	<frozen importlib._bootstrap>:870(_find_spec)
1	0.000	0.000	0.000	0.000	<frozen importlib._bootstrap_external>:743(get_code)
1	0.000	0.000	0.000	0.000	<frozen importlib._bootstrap_external>:1149(find_spec)
1	0.000	0.000	0.000	0.000	<frozen importlib._bootstrap_external>:1117(_get_spec)
3	0.000	0.000	0.000	0.000	<frozen importlib._bootstrap_external>:1233(find_spec)

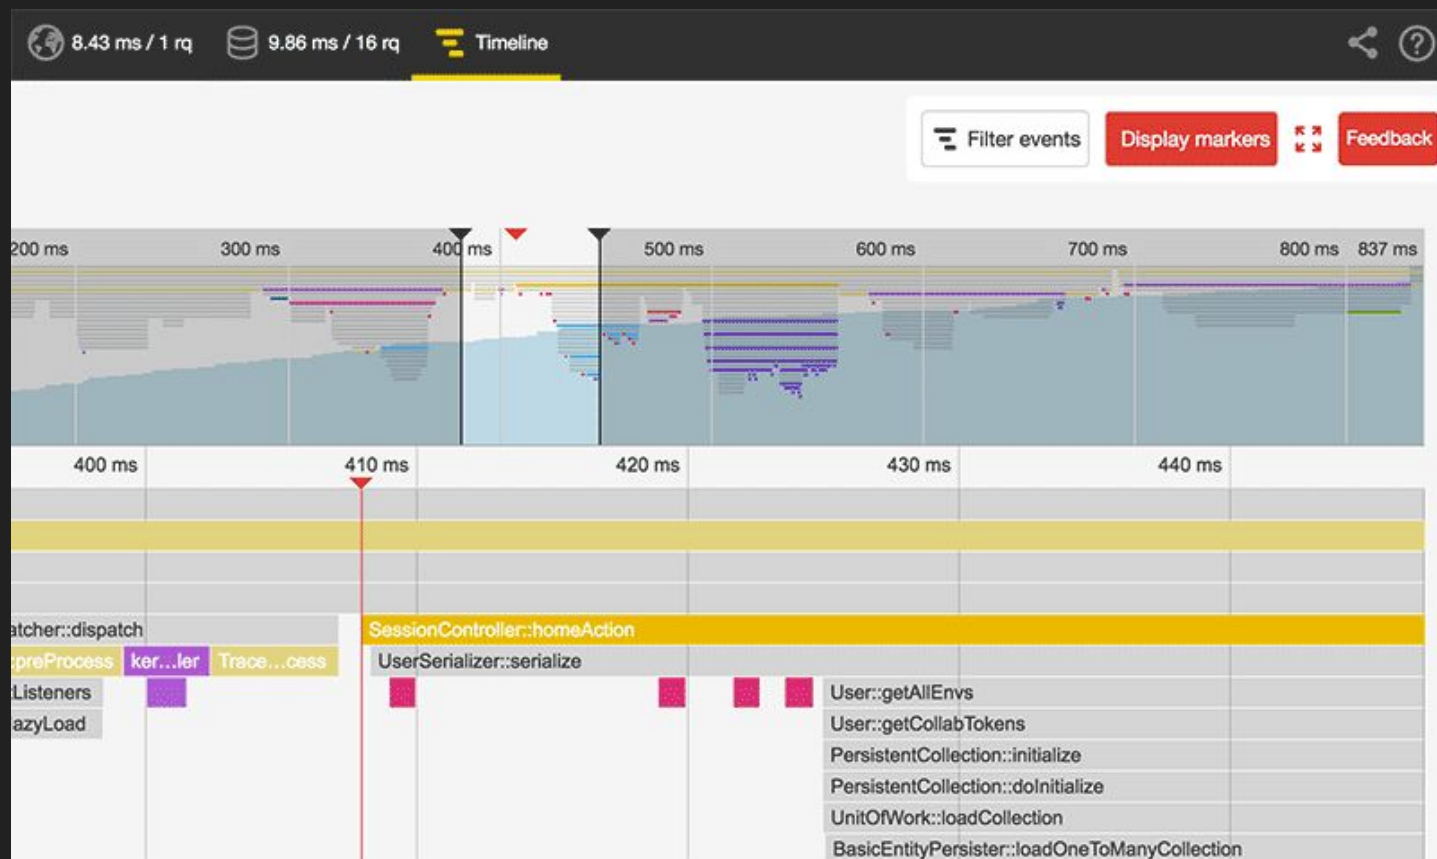
Callgraph



Flamegraph



Custom



Examples

line_profiler (line tracing)

- Measures per-line **wall time**.
- Can inspect its output from cmdline.

```
In [3]: load_ext line_profiler
```

```
In [4]: lprun -f get_books_by_library get_books_by_library()
```

```
Timer unit: 1e-06 s
```

```
Total time: 47.977 s
```

```
File: <ipython-input-1-5cb238008f5e>
```

```
Function: get_books_by_library at line 1
```

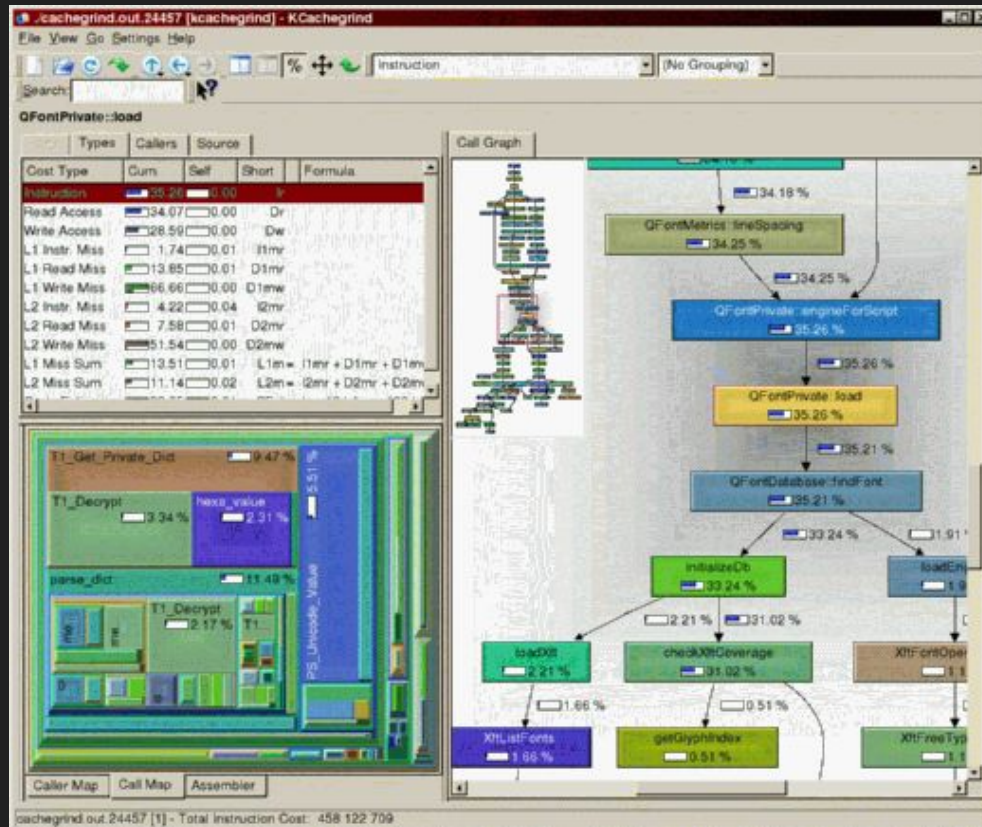
Line #	Hits	Time	Per Hit	% Time	Line Contents
1					def get_books_by_library():
2	1	312.0	<u>312.0</u>	0.0	libraries = Library.objects.all()
3	1	1.0	1.0	0.0	result = {}
4	1001	30298.0	30.3	0.1	for library in libraries:
5	1000	566487.0	566.5	1.2	books_in_library = library.book_set.all()
6	1000	47379897.0	47379.9	98.8	result[library.id] = list(books_in_library)
7					
8	1	0.0	0.0	0.0	return result

```
In [5]:
```

Yappi (function tracing)

- Measures per-function `wall/CPU time`.
- Can profile `multithreaded/async` applications and show `per-thread` traces.
- Supports profiling `asyncio` and `gevent(new)` applications.
- Can use `KCacheGrind` to visualize its output as a callgraph.

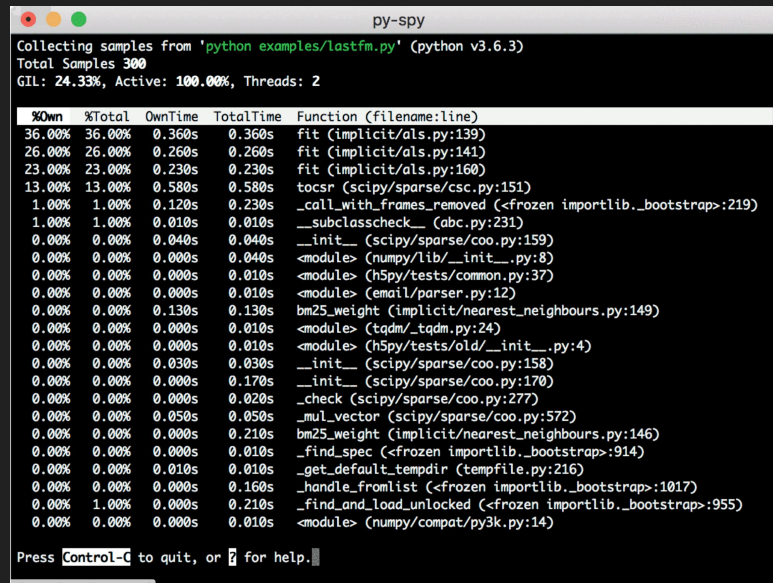

Yappi (function tracing)



py-spy (external sampling)

- Measures per-function CPU time.
- Minimal overhead.(~%10-%15)
- top like CLI, can output to FlameGraph or SpeedScope formats.
- Show GIL contention whenever possible.

py-spy (external sampling)



tracemalloc (memory profiler)

- Included in the stdlib. as of 3.4. (There is a backport for 2.x)
- Traces all `malloc/realloc/free` calls and saves a `traceback` along with the allocation.

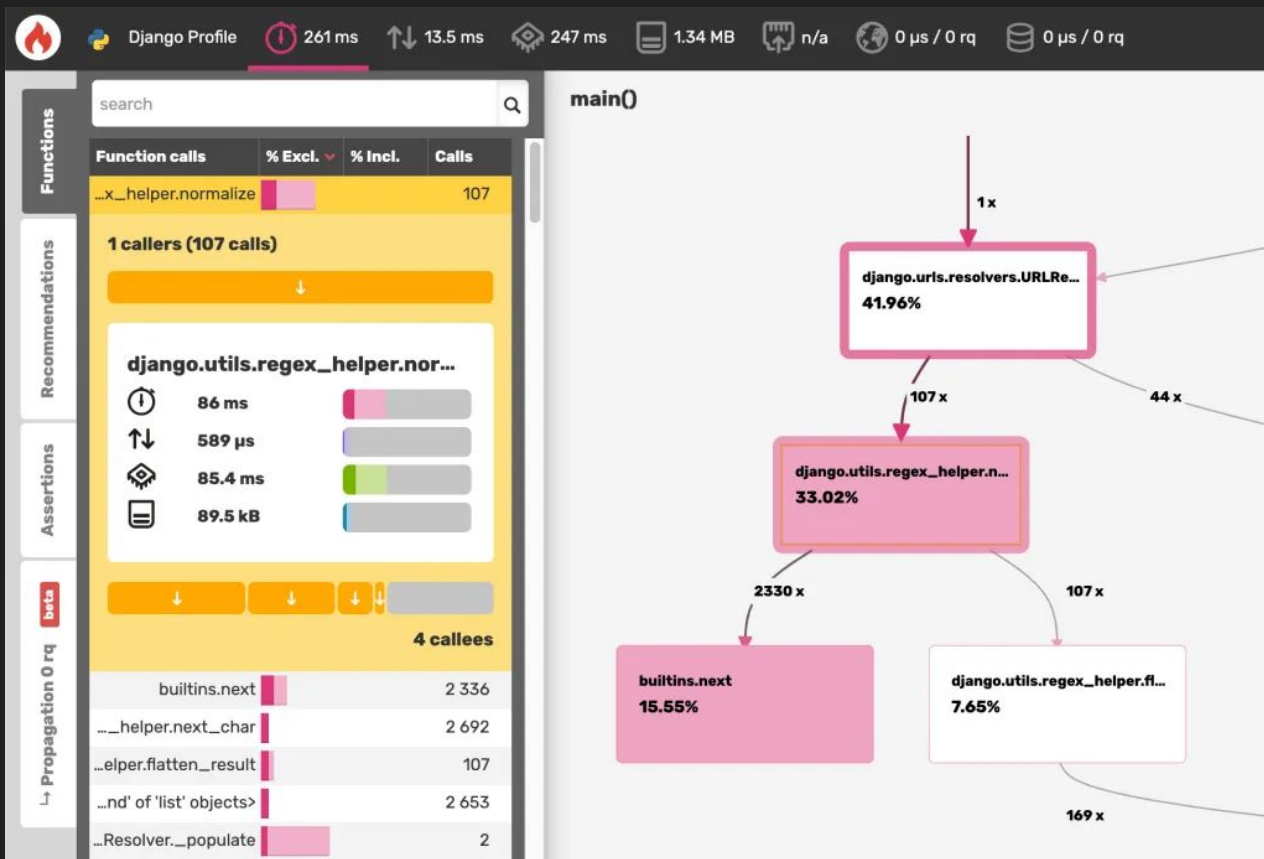
Example output:

```
tm.py:15: size=4656 B, count=1, average=4656 B
tm.py:14: size=64 B, count=1, average=64 B
tm.py:14: size=40 B, count=1, average=40 B
```

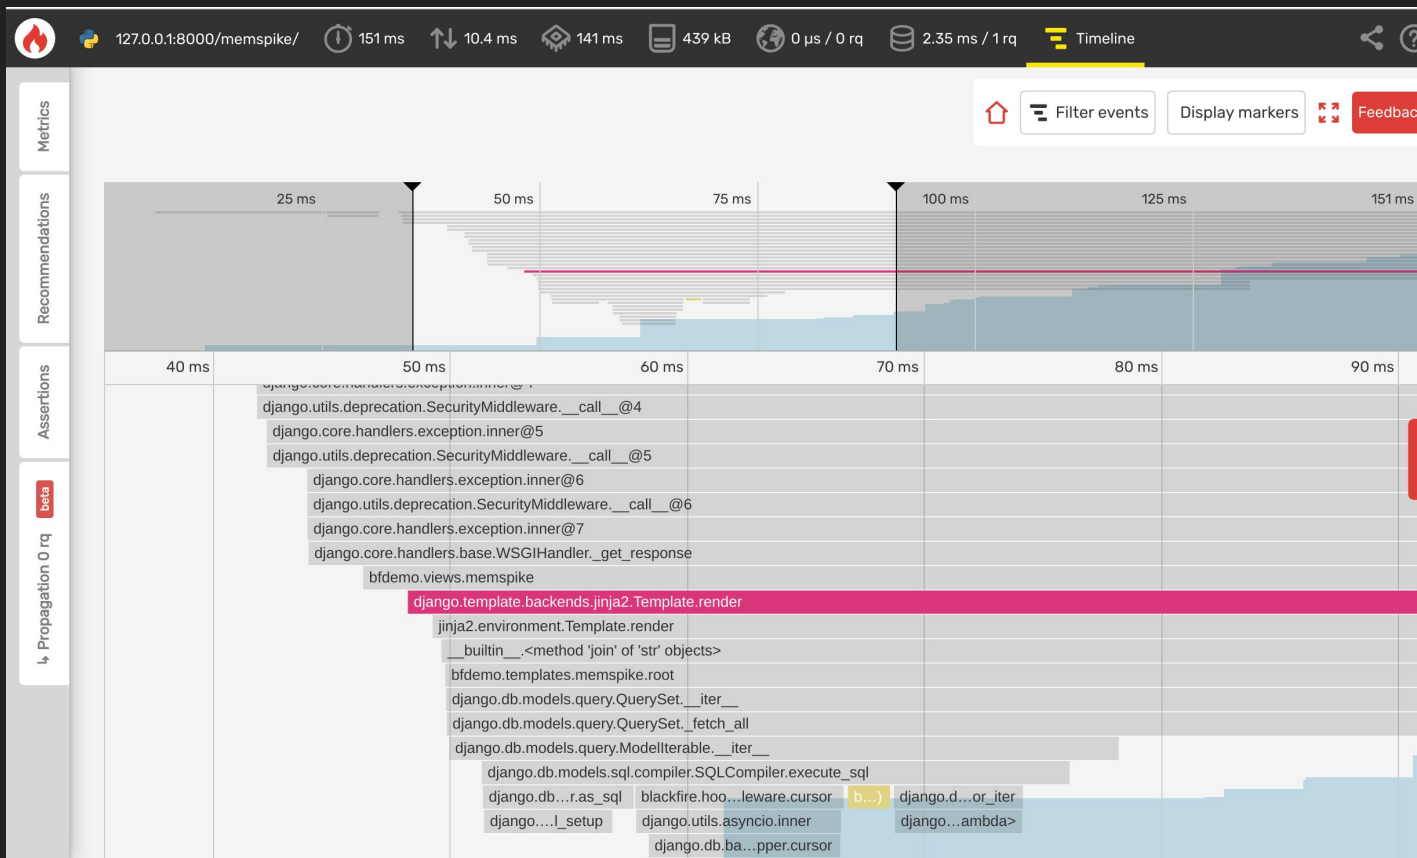
Blackfire (on-demand tracing)

- Measures per-function **wall/CPU time, memory** and some app. specific metrics.
- Enabled only **on-demand**.
- Profiles **microservice** requests.
- Has a web application to show its output as **callgraph** and **FlameGraph with time(Timeline)**.

Blackfire (on-demand tracing)



Blackfire (on-demand tracing)



Common Rules

Common Rules

- Always **measure**.
- Focus first on **Architecture/Design/Algorithms**.
- Know your **data** and how you access it(e.g: asymptotic complexity).
- Check if **standard library** has a function for your case.
- Avoid micro optimizations.
- If you really need performance, then there are tools. **Cython, Numba** or directly write code in C.

Thank you!

twitter.com/sumercip
github.com/sumerc
sumercip.com